

Smart Phone Sentiment Analysis and Predictions for Medical App Development

Project background: This project involved working with a government health organization. The overall objective was to develop medical apps to allow aid workers to manage health conditions by facilitating communication with medical professionals elsewhere. However, to limit cost and ensure consistency of training, the agency required the apps be developed for only one cell phone type.

My task was to perform web analysis to determine the prevalence of sentiment toward a short list of devices that were previously determined capable of executing the app's functions. The goal was to select one device by conducting a broad-based sentiment analysis.

The work flow for this project was divided into **two sections**. **The first involved** the identification and data collections. In the initial step of collecting the input web addresses the wet paths file for the month of interest was downloaded from the Common Crawl. Since the project calls for sentiment mining, work was focused on a subset of WARC files that only contain text:WET. Data mining was done Using Amazon's EMR with previously written python mapper and reducer scripts to access and compile web pages from the Common Crawl. An initial exploratory EMR Cluster on one web address was created and run to ensure the job would run smoothly. The data output was then downloaded and using the Anaconda prompt and a concatenate program the data was compiled into a single csv file. Once smooth work flow was confirmed, an additional 255 web addresses were searched by creating a large cluster using the AWS CLI. For this step a .bdf file was created containing the 255 addresses of interest, the addresses were mapped to the AWS S3:// commoncrawl/ folder, and the .bdf folder was then converted to a Json file, a cluster was created on the AWS CLI using Anaconda prompt and the Json file was run in the EMR cluster. The data output files (255 files) of the job was then down loaded and using a concatenate program combined into one csv folder containing over 17,000 observations and attributes describing sentiment toward listed cell phone types. This data set can them be used and predict sentiment to determine the preferred device for our medical app development.

The **Second section** of this work involved the development of machine learning models that predict cell phone sentiment for two preferred cell phones (iphone & Samsung Galaxy) as determined by the health government Agency. For training data sets we have been provided an iphone dataset and a galaxy dataset. These data sets include counts of pertinent words, describing sentiment, for approximately 12,000 observations. The assigned values in the sentiment attribute is representative of the overall sentiment toward the device. These values were assigned by an independent group by read each review rating sentiment. The rating scale is : "0 = very negative, 1 = negative, 2 = somewhat negative, 3 = somewhat positive, 4 = positive, 5 = very positive.

Prediction will be made on the data set that was collected for AWS in the first step of this project.

Begin by looking at iPhone training data set:

Set up for parallel processing for faster processing

```
install.packages("doParallel")
library(doParallel)
## find out how many cores are on your machine
detectCores()
## create cluster with desired number of cores. dont use them all.
cl <- makeCluster(2)
## register cluster
registerDoParallel(cl)
## confirm how many cores are now "assigned" to R and Rstudio
getDoParWorkers()
## stop Cluster. after performing tasks, stop your cluster
## use stopCluster(cl)
```

install packages

```
library(readr)
library(caret)
library(plotly)
library(corrplot)
library(dplyr)
```

#upload the data

```
iphone_smallmatrix <- read.csv("iphone_smallmatrix_labeled_8d.csv")
```

Data at a glance

```
names(iphone_smallmatrix)
```

```
[1] "iphone"           "samsunggalaxy"   "sonyxpria"
[4] "nokialumina"      "htcphone"        "ios"
[7] "googleandroid"    "iphonecampos"    "samsungcampos"
[10] "sonycampos"       "nokiacampos"     "htccampos"
[13] "iphonecamneg"     "samsungcamneg"   "sonycamneg"
[16] "nokiacamneg"      "htccamneg"       "iphonecamunc"
[19] "samsungcamunc"    "sonycamunc"      "nokiacamunc"
[22] "htccamunc"        "iphonedispos"    "samsungdispos"
[25] "sonydispos"       "nokiadispos"     "htcdispos"
[28] "iphonedisneg"     "samsungdisneg"   "sonydisneg"
[31] "nokiadisneg"      "htcdisneg"       "iphonedisunc"
[34] "samsungdisunc"    "sonydisunc"      "nokiadisunc"
[37] "htcdisunc"        "iphoneperpos"    "samsungperpos"
[40] "sonyperpos"       "nokiaperpos"     "htcperpos"
[43] "iphoneperneg"     "samsungperneg"   "sonyperneg"
[46] "nokiaperneg"      "htcperneg"       "iphoneperunc"
[49] "samsungperunc"    "sonyperunc"      "nokiaperunc"
[52] "htcperunc"        "iosperpos"       "googleperpos"
[55] "iosperneg"        "googleperneg"    "iosperunc"
[58] "googleperunc"     "iphonesentiment"
```

Explanation of the names:

- Iphone, samsunggalaxy, sonyxperia , nokialumina, and htcphone are the five phones on the list that have been determined to be able to perform the app's fuctions.
- ios is the count of mentions of the iphone operating system on the review
- iphonecampos- is the count mentions of positive sentiment towards the iphone camera
- iphonedisunc- is the count of unclear setiment metions toward the iphone display
- iphoneperneg- is the count mentions of negative sentiment towards the iphone performance

check data types of attributes

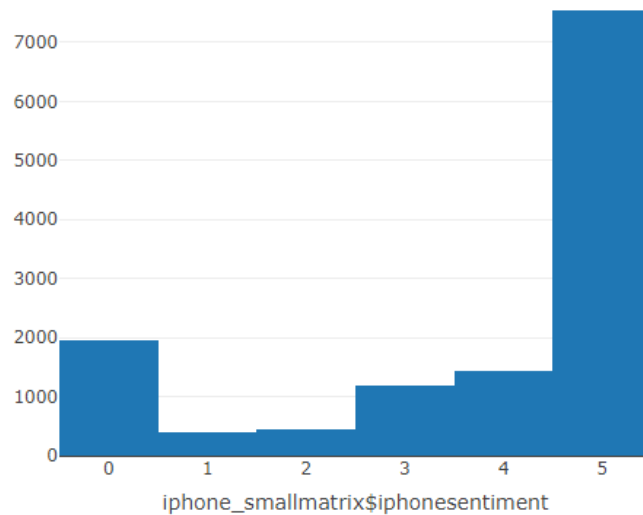
```
str(iphone_smallmatrix)
str(iphone_smallmatrix$iphonesentiment)
```

```
str(iphone_smallmatrix$iphonesentiment)
int [1:12973] 0 0 0 0 0 4 4 0 0 0 ...
```

upon checking the structure of the attributes it is clear that all attributes are integers. For visualization I have only included the structure of the iphonesentiment, which will also be the dependent variable for machine learning and sentiment prediction.

to further understand the data, I then plot & study the distribution of the #dependent variable scores.

```
plot_ly(iphone_smallmatrix, x= ~iphone_smallmatrix$iphonesentiment, type= 'histogram')
```



Based on the ditrubution plot it is clear that over 50% of the reviews have a positive sentiment toward the iphone.

check for missing data and adress if needed

```
sum(is.na(iphone_smallmatrix))
```

```
sum(is.na(iphone_smallmatrix))
[1] 0
```

No missing data was found in our data set.

For feature selection we will look at correlation, near zero variance, and recursive feature elimination. For each approach we will create a new data set to test our machine learning models on to determine The best one.

Check for near Zero Variance: Zero or zero variance in variables shows a constant or near constant Predictors across all the observations. This type of constant predictor does not provide much information and may also bog down some of the machine learning algorithms. Attributes with Zero or near Zero variance will be removed.

study correlation of between features and dependent variable

4

```
# based on correlation data no high correlation thus no features will be removed
# create a new data frame from the iphone_smallmatrix data frame.
# in this case, this new data frame retains all the feature for "out of the box modeling"
```

```
iphoneCor <- iphone_smallmatrix
# to remove a feature if it were needed we would use
# iphoneCor$featuretoremove <- Null
```

```
# search and remove near zero variance features. create new data set without nearzero variance
# using nearZeroVar with saveMetrics =True returns an object containing
# at a table: with frequency ration, % unique, zero variance and
# near zero variance. it allow us to see the attributes with zero variance.
```

```
nzvMetrics <- nearZeroVar(iphone_smallmatrix, saveMetrics = TRUE)
nzvMetrics
```

	freqRatio	percentUnique	zeroVar	nzv
iphone	5.041322	0.20812457	FALSE	FALSE
samsunggalaxy	14.127336	0.05395822	FALSE	FALSE
sonyxpria	44.170732	0.03854159	FALSE	TRUE
nokialumina	497.884615	0.02312495	FALSE	TRUE
htcphone	11.439614	0.06937486	FALSE	FALSE
ios	27.735294	0.04624990	FALSE	TRUE
googleandroid	61.247573	0.04624990	FALSE	TRUE
iphonecampos	10.524697	0.23124952	FALSE	FALSE
samsungcampos	93.625000	0.08479149	FALSE	TRUE
sonycampos	348.729730	0.05395822	FALSE	TRUE
nokiacampos	1850.142857	0.08479149	FALSE	TRUE
htccampos	79.272152	0.16958298	FALSE	TRUE
iphonecamneg	19.517529	0.13104139	FALSE	TRUE
samsungcamneg	100.132812	0.06937486	FALSE	TRUE
sonycamneg	1851.285714	0.04624990	FALSE	TRUE
nokiacamneg	2158.833333	0.06166654	FALSE	TRUE
htccamneg	93.444444	0.11562476	FALSE	TRUE
iphonecamunc	16.764205	0.16187466	FALSE	FALSE
samsungcamunc	74.308140	0.06937486	FALSE	TRUE
sonycamunc	588.318182	0.03854159	FALSE	TRUE
nokiacamunc	2591.200000	0.05395822	FALSE	TRUE
htccamunc	50.548000	0.12333308	FALSE	TRUE
iphonedispos	6.792440	0.24666615	FALSE	FALSE
samsungdispos	97.061069	0.13104139	FALSE	TRUE
sonydispos	331.076923	0.06937486	FALSE	TRUE
nokiadispos	1438.777778	0.09249981	FALSE	TRUE
htcdispos	64.694301	0.20041625	FALSE	TRUE
iphonedisneg	10.084428	0.18499961	FALSE	FALSE
samsungdisneg	99.155039	0.10791644	FALSE	TRUE
sonydisneg	2159.333333	0.06937486	FALSE	TRUE
nokiadisneg	1850.142857	0.08479149	FALSE	TRUE
htcdisneg	88.492958	0.14645803	FALSE	TRUE
iphonedisunc	11.471875	0.20812457	FALSE	FALSE
samsungdisunc	74.255814	0.09249981	FALSE	TRUE
sonydisunc	719.222222	0.05395822	FALSE	TRUE
nokiadisunc	1619.375000	0.04624990	FALSE	TRUE
htcdisunc	50.590361	0.13874971	FALSE	TRUE
iphoneperpos	9.297834	0.19270793	FALSE	FALSE
samsungperpos	94.200000	0.10791644	FALSE	TRUE

sonyperpos	416.870968	0.06166654	FALSE	TRUE
nokiaperpos	2158.000000	0.08479149	FALSE	TRUE
htcperpos	74.279762	0.19270793	FALSE	TRUE
iphoneperneg	11.054137	0.16958298	FALSE	FALSE
samsungperneg	101.650794	0.10020812	FALSE	TRUE
sonyperneg	2159.666667	0.07708317	FALSE	TRUE
nokiaperneg	3237.250000	0.09249981	FALSE	TRUE
htcperneg	94.428571	0.15416635	FALSE	TRUE
iphoneperunc	13.018349	0.12333308	FALSE	FALSE
samsungperunc	86.500000	0.09249981	FALSE	TRUE
sonyperunc	3240.250000	0.04624990	FALSE	TRUE
nokiaperunc	1850.428571	0.06937486	FALSE	TRUE
htcperunc	50.055556	0.15416635	FALSE	TRUE
iosperpos	153.373494	0.09249981	FALSE	TRUE
googleperpos	98.592308	0.06937486	FALSE	TRUE
iosperneg	141.744444	0.09249981	FALSE	TRUE
googleperneg	99.403101	0.08479149	FALSE	TRUE
iosperunc	135.893617	0.07708317	FALSE	TRUE
googleperunc	96.443609	0.07708317	FALSE	TRUE
iphonesentiment	3.843017	0.04624990	FALSE	FALSE

**# nearzeroVar () with saveMetrics = FALSE retruns a vector which can then
be used to easily and quckily remove the zero variance features.**

```
nvz <- nearZeroVar(iphone_smallmatrix, saveMetrics = FALSE)
nvz
```

```
[1] 3 4 6 7 9 10 11 12 13 14 15 16 17 19 20 21 22 24 25 26 27 29 30
[24] 31 32 34 35 36 37 39 40 41 42 44 45 46 47 49 50 51 52 53 54 55 56 57
[47] 58
```

remove nearzero variace features and create new dataset

```
iphoneNZV <- iphone_smallmatrix[,-nvz]
str(iphoneNZV)
```

Recursive Feature Elimination:

Recursive Feature Elimination: Here I Apply random forest functions to build the RFE model,
and cross validation to avoid overfitting

sample data before using RFE.

```
set.seed(123)
iphoneSample <- iphone_smallmatrix[sample(1:nrow(iphone_smallmatrix), 1000,replace = FALSE),]
```

set up rfeControl with randomforest functions, repeated cross validation and no updates

```
ctrl <- rfeControl(functions = rfFuncs,
  method = "repeatedcv",
  repeats = 5,
  verbose = FALSE)
```

Use rfe sampled data and omit the dependent variable (attribute 59, iphonesentiment)

```
rfeResults <- rfe(iphoneSample[,1:58],
  iphoneSample$iphonesentiment,
  sizes = (1:58), rfeControl = ctrl)
```

#get the results

rfeResults

Recursive feature selection

Outer resampling method: Cross-validated (10 fold, repeated 5 times)

Resampling performance over subset size:


Variables	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD	Selected
1	1.465	0.3125	1.1029	0.10778	0.10233	0.07811	
2	1.468	0.3209	1.1396	0.09954	0.10063	0.06861	
3	1.480	0.3165	1.1582	0.09438	0.09644	0.06490	
4	1.484	0.3208	1.1641	0.09356	0.09943	0.06499	
5	1.456	0.3544	1.1438	0.09178	0.10008	0.06269	
6	1.380	0.3914	1.0207	0.10184	0.09793	0.06701	
7	1.368	0.4023	1.0047	0.10361	0.09790	0.06596	
8	1.359	0.4108	0.9989	0.10481	0.09849	0.06527	
9	1.336	0.4275	0.9427	0.11560	0.10242	0.07301	
10	1.334	0.4285	0.9418	0.11219	0.09944	0.06999	
11	1.334	0.4294	0.9432	0.11624	0.10228	0.07394	
12	1.337	0.4266	0.9236	0.11800	0.10212	0.07682	
13	1.335	0.4280	0.9265	0.11851	0.10235	0.07546	
14	1.334	0.4289	0.9288	0.11791	0.10279	0.07622	
15	1.332	0.4310	0.9146	0.11959	0.10264	0.07770	
16	1.331	0.4318	0.9175	0.11699	0.10140	0.07650	
17	1.327	0.4346	0.9201	0.11624	0.09956	0.07526	
18	1.326	0.4356	0.9095	0.11788	0.10101	0.07706	
19	1.326	0.4357	0.9141	0.11725	0.10001	0.07633	*
20	1.326	0.4354	0.9171	0.11677	0.09974	0.07571	
21	1.327	0.4353	0.9100	0.11688	0.09934	0.07669	
22	1.328	0.4344	0.9139	0.11642	0.09921	0.07580	
23	1.328	0.4341	0.9174	0.11577	0.09896	0.07544	
24	1.329	0.4338	0.9115	0.11774	0.10038	0.07694	
25	1.328	0.4344	0.9129	0.11655	0.10000	0.07678	
26	1.327	0.4350	0.9152	0.11598	0.09915	0.07591	
27	1.327	0.4348	0.9094	0.11732	0.10013	0.07667	
28	1.328	0.4345	0.9118	0.11632	0.09931	0.07606	
29	1.328	0.4341	0.9144	0.11565	0.09912	0.07616	
30	1.328	0.4344	0.9098	0.11662	0.09968	0.07680	
31	1.328	0.4344	0.9119	0.11642	0.09978	0.07640	
32	1.328	0.4341	0.9145	0.11681	0.10004	0.07709	
33	1.328	0.4344	0.9096	0.11728	0.10005	0.07772	
34	1.327	0.4350	0.9104	0.11631	0.09928	0.07590	
35	1.327	0.4350	0.9128	0.11582	0.09919	0.07652	
36	1.328	0.4347	0.9092	0.11735	0.09993	0.07797	
37	1.328	0.4346	0.9111	0.11719	0.09974	0.07706	
38	1.328	0.4341	0.9129	0.11686	0.10025	0.07747	
39	1.327	0.4349	0.9090	0.11735	0.10007	0.07719	
40	1.327	0.4347	0.9104	0.11585	0.09926	0.07656	
41	1.327	0.4347	0.9126	0.11736	0.10008	0.07733	
42	1.328	0.4347	0.9086	0.11667	0.09928	0.07725	
43	1.328	0.4342	0.9104	0.11637	0.09966	0.07712	
44	1.327	0.4347	0.9114	0.11646	0.09971	0.07777	
45	1.327	0.4354	0.9081	0.11769	0.09995	0.07786	
46	1.328	0.4345	0.9103	0.11682	0.09984	0.07748	
47	1.328	0.4345	0.9118	0.11698	0.09998	0.07785	
48	1.328	0.4347	0.9090	0.11713	0.09983	0.07756	
49	1.328	0.4345	0.9104	0.11636	0.09945	0.07731	
50	1.328	0.4345	0.9121	0.11695	0.09978	0.07763	
51	1.328	0.4341	0.9102	0.11676	0.09943	0.07797	
52	1.328	0.4345	0.9105	0.11740	0.10010	0.07803	
53	1.328	0.4343	0.9122	0.11676	0.09975	0.07714	
54	1.328	0.4344	0.9091	0.11704	0.09981	0.07735	
55	1.328	0.4343	0.9108	0.11700	0.09965	0.07778	
56	1.328	0.4341	0.9112	0.11811	0.10062	0.07889	
57	1.329	0.4332	0.9090	0.11745	0.09983	0.07946	
58	1.329	0.4336	0.9088	0.11771	0.10066	0.07903	

The top 5 variables (out of 19):

iphone, googleandroid, iphonedispos, iphonedisneg, samsunggalaxy

Plot rfeResults results

```
plot(rfeResults, type=c("g","o"))
```



The resulting table and plot display each subset and its statistical values. An asterisk denotes the number of features that is judged the most optimal from RFE. Both the table and the plot show the optimal subset is the one that has the lowest RMSE. In this case it consists of 19 variables.

create a new data set with the rfe recommended features

```
iphoneRFE <- iphone_smallmatrix[,predictors(rfeResults)]
```

add the dependent variable to iphoneRFE

```
iphoneRFE$iphonesentiment <- iphone_smallmatrix$iphonesentiment  
str(iphoneRFE)
```

Preprocessing:

Since this is a classification problem our dependent variable will be converted into a factor. This will be done to all the data sets created via feature selection and the original data set. (in this case, no features were removed from the original data set due to correlation. Thus, since `iphone_smallmatrix` and `iphoncorr` data set are the same only one will be used for modeling.

```
iphone_smallmatrix$iphonesentiment <- as.factor(iphone_smallmatrix$iphonesentiment)  
iphoneCor$iphonesentiment <- as.factor(iphoneCor$iphonesentiment)  
iphoneNZV$iphonesentiment <- as.factor(iphoneNZV$iphonesentiment)  
iphoneRFE$iphonesentiment <- as.factor(iphoneRFE$iphonesentiment)
```

Model Development and Evaluation

For Model Development in this project now data sampling will be used. The small matrix data set was classified by hand, thus all observations will be used for training and testing. Four classification machine learning algorithms (C5.0, Random Forest, KNN, and SVM) using the original data set containing all of the features. Since the goal is to find the best combination of algorithm and feature selection method for predicting sentiment, the best model based on accuracy and Kappa will be selected to model on all

feature selection data sets. Kappa and accuracy parameters will be compared and the best combination will be chosen to predict sentiment on iPhone in the large data set prepared in step one of this project.

create data partition and training and testing sets

```
set.seed(123)
iphoneinTrain <- createDataPartition( y = iphone_smallmatrix$iphonesentiment,p=0.70, list = FALSE)
str(iphoneinTrain)
iphonettraining <- iphone_smallmatrix [iphoneinTrain,]
iphonetesting <- iphone_smallmatrix [-iphoneinTrain,]
nrow(iphonettraining)
nrow(iphonetesting)
```

#create fit control

```
iphonefitControl <- trainControl(method= "repeatedcv", number=10,repeats = 1)
```

create C5.0 model for classification iphone_smallmatrix data frame

```
library(C50)
iphone_smallmatrix_C5.0model<- train (iphonesentiment~.,data = iphonetraining, method = "C5.0",
trControl=iphonefitControl)
iphone_smallmatrix_C5.0model
varImp(iphone_smallmatrix_C5.0model)
# predicting from iphone testing
iphone_small_C5.0_predictions <- predict(iphone_smallmatrix_C5.0model, iphonetesting)
```

create random forest model for classification iphone_smallmatrix

```
iphone_smallmatrix_RFmodel<- train (iphonesentiment~.,data = iphonetraining, method = "rf",
trControl=iphonefitControl)
iphone_smallmatrix_RFmodel
# predicting from iphone testing
iphone_small_RF_predictions <- predict(iphone_smallmatrix_RFmodel, iphonetesting)
```

#create kkNN model for classification of iphone_smallmatrix

```
iphone_smallmatrix_KKNNmodel<- train (iphonesentiment~.,data = iphonetraining, method = "knn",
trControl=iphonefitControl)
iphone_smallmatrix_KKNNmodel
# predicting from iphone testing
iphone_small_KKNN_predictions <- predict(iphone_smallmatrix_KKNNmodel, iphonetesting)
```

create SVM model for classification of iphone_smallmatrix data frame

```
library(e1071)
iphone_smallmatrix_SVMmodel<- train (iphonesentiment~.,data = iphonetraining, method =
"svmLinear2", trControl=iphonefitControl)
iphone_smallmatrix_SVMmodel
# predicting from iphone testing
iphone_small_SVM_predictions <- predict(iphone_smallmatrix_SVMmodel, iphonetesting)
```

Evaluating the models

To evaluate the models I first use resamples to determine and obtain model metrics from the training (iphonettraining) data partition. Second I will use Postresample to compare the prediction to the groundtruth in the testing (iphonetesting) data partition.

Accuracy describes the number of instances that were classified correctly. kappa describes a comparison between an observed accuracy and an expected accuracy. Expected Accuracy is the accuracy that any random classifier would be expected to achieve. Expected Accuracy is directly related to the number of instances of each class combined with the number of instances that the machine learning classifier agreed with to be ground truth. Overall, Kappa Score is less misleading than simply using accuracy.

##Obtaining accuracy and kappa summary of the for iphone_small matrix models using (resamples)

```
iphone_smallmatrix_ModelData <- resamples(list(C50 = iphone_smallmatrix_C5.0model, RF =  
iphone_smallmatrix_RFmodel, kknn = iphone_smallmatrix_KKNNmodel, SVM2 =  
iphone_smallmatrix_SVMmodel))  
summary(iphone_smallmatrix_ModelData)
```

```
Call:  
summary.resamples(object = iphone_smallmatrix_ModelData)
```

```
Models: C50, RF, kknn, SVM2  
Number of resamples: 10
```

Accuracy	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
C50	0.7577093	0.7663268	0.7713034	0.7715533	0.7778702	0.7850055	0
RF	0.7546755	0.7686389	0.7744772	0.7731025	0.7769824	0.7913907	0
kknn	0.3171806	0.3275526	0.3309512	0.3313906	0.3382941	0.3428886	0
SVM2	0.6938326	0.6979371	0.7077627	0.7079155	0.7155428	0.7282728	0

Kappa	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
C50	0.5241782	0.5422098	0.5578319	0.5561141	0.5685017	0.5863730	0
RF	0.5257395	0.5535973	0.5658198	0.5627197	0.5693608	0.6015349	0
kknn	0.1395009	0.1586956	0.1626841	0.1621276	0.1696927	0.1761262	0
SVM2	0.3723279	0.3848960	0.4056711	0.4077320	0.4245260	0.4648812	0

From using "resamples" function we can see that the model with the highest accuracy and kappa values, during model training, is the random forest algorithm. However, the C5.0 model is also with very similar values.

To study which model is best suited for predictions we will test the predictions against the ground truth on the testing data partition (iphonetesting) we created. This is done using the “postresample” function.

**# evaluate models using postresample to compare predictions to ground truth
in testing (iphonetesting) datapartition.**

#evaluate C5.0

postResample(iphone_small_C5.0_predictions, iphonetesting\$iphonesentiment)

#evaluate RF


postResample(iphone_small_RF_predictions, iphonetesting\$iphonesentiment)

#evaluate KNN

postResample(iphone_small_KNN_predictions, iphonetesting\$iphonesentiment)

#evaluate SVM

postResample(iphone_small_SVM_predictions, iphonetesting\$iphonesentiment)



Upon evaluation of the predictions and ground truth values we can see that model that performed the best are the C5.0 and Random Forest algorithms. For time saving interest I will use the C5.0 machine learning algorithm to test and determine the best feature selection strategy for our classification sentiment prediction strategy.

The C5.0 model has been selected to be used on the feature selection data sets. This model has already been run on the iPhone small matrix data set. So I will only train it using the NZV and RFE feature selection data sets.

#iphoneNZV feature selection data set

```
iphone_NZVInTrain <- createDataPartition( y = iphoneNZV$iphonesentiment, p=0.70, list = FALSE)
str(iphone_NZVInTrain)
iphone_NZVtraining <- iphoneNZV [iphone_NZVInTrain,]
iphone_NZVtesting <- iphoneNZV [-iphone_NZVInTrain,]
nrow(iphone_NZVtraining)
nrow(iphone_NZVtesting)
```

#create fit control

```
iphone_NZVfitControl <- trainControl(method= "repeatedcv", number=10, repeats = 1)
```

create C5.0 model for classification iphone_smallmatrix data frame

```
library(C50)
iphone_NZV_C5.0model <- train (iphonesentiment~., data = iphone_NZVtraining, method = "C5.0",
trControl=iphone_NZVfitControl)
iphone_NZV_C5.0model
varImp(iphone_NZV_C5.0model)
# predicting from iPhone testing
iphone_NZV_C5.0_predictions <- predict(iphone_NZV_C5.0model, iphone_NZVtesting)
```

#iphoneRFE feature selection data set

```
iphone_RFEInTrain <- createDataPartition( y = iphoneRFE$iphonesentiment, p=0.70, list = FALSE)
str(iphone_RFEInTrain)
iphone_RFEtraining <- iphoneRFE [iphone_RFEInTrain,]
iphone_RFEtesting <- iphoneRFE [-iphone_RFEInTrain,]
nrow(iphone_RFEtraining)
nrow(iphone_RFEtesting)
```

#create fit control

```
iphone_RFEfitControl <- trainControl(method= "repeatedcv", number=10, repeats = 1)
```

create C5.0 model for classification iphone_smallmatrix data frame

```
iphone_RFE_C5.0model <- train (iphonesentiment~., data = iphone_RFEtraining, method = "C5.0",
trControl=iphone_RFEfitControl)
iphone_RFE_C5.0model
varImp(iphone_RFE_C5.0model)
# predicting from iPhone testing
iphone_RFE_C5.0_predictions <- predict(iphone_RFE_C5.0model, iphone_RFEtesting)
```

#evaluate C5.0 models with different feature selection data sets based on the training data set using resamples

##Obtaining accuracy and kappa summary of the for C5.0 models using (resamples)

```
iphone_C5.0_ModelData <- resamples(list(C50_is = iphone_smallmatrix_C5.0model, C5.0_NZV =
iphone_NZV_C5.0model, C5.0_RFE = iphone_RFE_C5.0model))
summary(iphone_C5.0_ModelData)
```

```
Call:
summary.resamples(object = iphone_C5.0_ModelData)
```

Models: C50_is, C5.0_NZV, C5.0_RFE
Number of resamples: 10

Accuracy

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
C50_is	0.7577093	0.7663268	0.7713034	0.7715533	0.7778702	0.7850055	0
C5.0_NZV	0.7431092	0.7529578	0.7557756	0.7572416	0.7613672	0.7750827	0
C5.0_RFE	0.7588106	0.7685757	0.7763073	0.7740807	0.7798843	0.7843784	0

Kappa

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
C50_is	0.5241782	0.5422098	0.5578319	0.5561141	0.5685017	0.5863730	0
C5.0_NZV	0.4937573	0.5112421	0.5204902	0.5222698	0.5315565	0.5598355	0
C5.0_RFE	0.5278570	0.5512829	0.5660371	0.5620452	0.5735504	0.5841558	0

Evaluate C5.0 models with different feature selection approaches by using postresample

#evaluate C5.0 with iphone_smallmatrix data set


```
postResample(iphone_small_C5.0_predictions, iphonetesting$iphonesentiment)
```

#evaluate C5.0 with iphone_NZV data set

```
postResample(iphone_NZV_C5.0_predictions, iphonetesting$iphonesentiment)
```

#evaluate C5.0 with iphoneRFE data set

```
postResample(iphone_RFE_C5.0_predictions, iphonetesting$iphonesentiment)
```



As we can see from the Chart both the no feature selection and the RFE feature selection approaches are compatible and out performed the nonzero variance approach. However, the accuracy and Kappa values are still not optimal and improvement is needed before we can confidently apply these models to our data set.

Model performance improvement by Performing Feature Engineering.

To improve the C5.0 model we change the number of levels of the dependent variable. From 6 to 4. Thus, level (0) very negative will be combined with level(1) negative and level (6) very positive will be combined with level (5) positive in a new data frame and the model will be tested to see if accuracy and kappa have improved.

```
# feature engineering to improve C5.0 model performance
# combine levels 0 and 1 into 1, and 5 and 6 into 5 using dplyr's recode function
#create a new dataset that will be used for recoding sentiment
iphone_Recode <- iphone_smallmatrix
# recode sentiment to combine factor levels 0 & 1 and 4 & 5
iphone_Recode$iphonesentiment <- recode(iphone_Recode$iphonesentiment, '0' = 1, '1' = 1, '2' = 2, '3'
= 3, '4' = 4, '5' = 4)
# inspect results
summary(iphone_Recode)
str(iphone_Recode)
# make iphonesentiment a factor
iphone_Recode$iphonesentiment <- as.factor(iphone_Recode$iphonesentiment)

# create the C5.0 model on the new data set iphone_Recode


# create data partition and training and testing sets
set.seed(123)
iphone_RC_inTrain <- createDataPartition( y = iphone_Recode$iphonesentiment,p=0.70, list = FALSE)
str(iphone_RC_inTrain)
iphone_RC_training <- iphone_Recode [iphone_RC_inTrain,]
iphone_RC_testing <- iphone_Recode [-iphone_RC_inTrain,]
nrow(iphone_RC_training)
nrow(iphone_RC_testing)

#create fit control
iphone_RC_fitControl <- trainControl(method= "repeatedcv", number=10, repeats = 1)

## create C5.0 model for classification iphone_smallmatrix data frame
library(C50)
iphone_RC_C5.0model<- train (iphonesentiment~.,data = iphone_RC_training, method = "C5.0",
trControl=iphone_RC_fitControl)
iphone_RC_C5.0model
varImp(iphone_RC_C5.0model)
# predicting from iphone testing
iphone_RC_C5.0_predictions <- predict(iphone_RC_C5.0model, iphone_RC_testing)

# obtain accuracy and kappa metrics for iphone_RC_C5.0 model to determine performance
postResample(iphone_RC_C5.0_predictions, iphone_RC_testing$iphonesentiment)
```

Model Name	Accuracy	Kappa
Iphone_RC_5.0model	0.85	0.63



The chart show that application of fearture engineering by using dplyr's recode function improve the C5.0 model performance for our data. Here Accuracy increased to 0.85 and kappa increased to 0.63. This values represent more accurate model that can be used for predicting iphone sentiment.

Galaxy Smart Phone:

The data for the second smart phone to consider (galaxy smart phone) is on a separate data set that was also hand classified for sentiment. Thus, this second data set was considered with the same feature selection, preprocessing, model testing, and feature engineering to optimize the selected model as the iPhone data set. The result is that the C5.0 model, with the feature engineering reducing number of levels for the dependent variable, also gave the best result. This report shows the modeling steps for only the optimal model. See chart below for performance comparison for the between the iPhone and galaxy data set models.

upload galaxy small matrix data set

```
galaxy_smallmatrix <- read.csv("galaxy_smallmatrix_labeled_9d.csv")
```

apply feature engineering with dplyr's recode method

create new data set that is to be used for recoding the levels of sentiment

```
galaxyRecode <- galaxy_smallmatrix
```

#recode the galaxy sentiment attribute levels 0&1 and 4&5

```
galaxyRecode$galaxysentiment <- recode(galaxyRecode$galaxysentiment, '0' = 1, '1' = 1, '2' = 2, '3' = 3, '4' = 4, '5' = 4)
```

inspect results

```
summary(galaxyRecode)
```

```
str(galaxyRecode)
```

make galaxy sentiment a factor

```
galaxyRecode$galaxysentiment <- as.factor(galaxyRecode$galaxysentiment)
```

prepare to train C5.0 model on galaxyRecode data set

create training and testing data sets from galaxyrecode data set

```
galaxyRC_inTrain <- createDataPartition(y = galaxyRecode$galaxysentiment,p=0.7, list = FALSE)
galaxyRC_training <- galaxyRecode[galaxyRC_inTrain,]
galaxyRC_testing <- galaxyRecode [-galaxyRC_inTrain,]
nrow(galaxyRC_training)
nrow(galaxyRC_testing)
```

#create fit control

```
galaxyRC_fitControl <- trainControl(method= "repeatedcv", number=10,repeats = 1)
```

apply and train the C5.0 model to the galaxyRecode data set

create C5.0 model for classification galaxyRecode data frame

```
library(C50)
galaxyRC_C5.0model<- train (galaxysentiment~.,data = galaxyRC_training, method = "C5.0",
trControl=galaxyRC_fitControl)
galaxyRC_C5.0model
varImp(galaxyRC_C5.0model)
```


predicting from iphone testing

```
galaxyRC_C5.0_predictions <- predict(galaxyRC_C5.0model, galaxyRC_testing)
```

obtain accuracy and kappa metrics for galaxyRC_C5.0 model to determine performance

```
postResample(galaxyRC_C5.0_predictions, galaxyRC_testing$galaxysentiment)
```

Model Name	Accuracy	Kappa
galaxy_RC_5.0model	0.84	0.60



Obtaining other model performance metrics for selected models

obtain a confusion matrix and model statistics using the confusionmatrix function

iphone_RC_C5.0_model

```
matrix_iphone_RC_C5.0 <- confusionMatrix(iphone_RC_C5.0_predictions,  
iphone_RC_testing$iphonesentiment)  
matrix_iphone_RC_C5.0
```

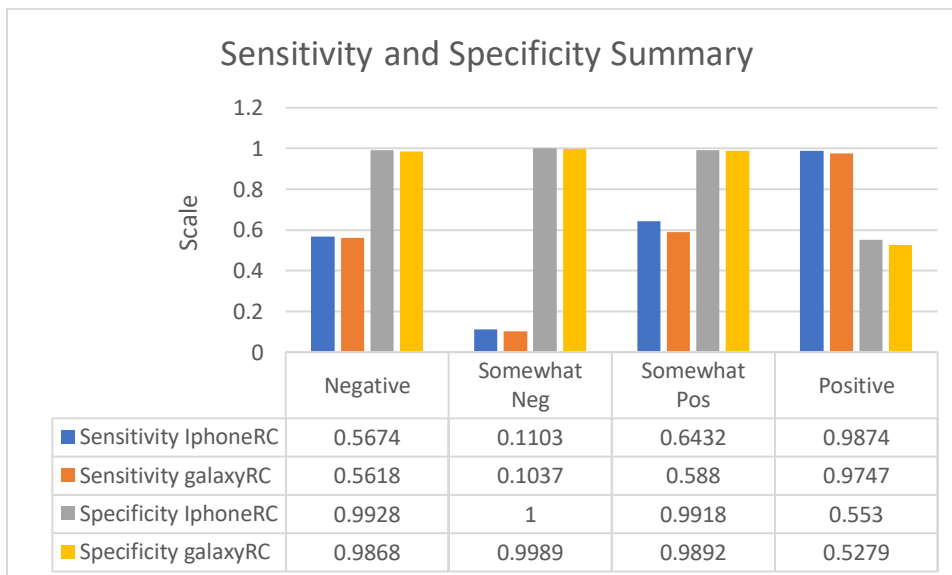
galaxyRC_C5.0_model

```
matrix_galaxyRC_C5.0 <- confusionMatrix(galaxyRC_C5.0_predictions,  
galaxyRC_testing$galaxysentiment)  
matrix_galaxyRC_C5.0
```

Summary of statistics for the models that will be used to classify the sentiment for the two selected smart phones

Classification model name	Accuracy	Kappa	95% CI	P value	Balance Accuracy	
					Neg. cat.	Pos. cat.
iPhone_RC_C5.0_model	0.85	0.63	0.8375, 0.8602	<2.2e-16	0.7801	0.7702
galaxyRC_C5.0_modle	0.84	0.6	0.8309, 0.8541	<2.2e-16	0.7843	0.7513

Summary of sensitivity and specificity statistics for the models that will be used to classify sentiment for the two selected smart phones.



Both Sensitivity and Specificity are important parameters in machine learning algorithms. Sensitivity describes the proportion of actual correct positive predictions while specificity describes the proportion of actual correct negative predictions.

Work with Large Data Matrix Collected via Amazon EMR to Make Prediction on iPhone and galaxy Sentiment

iphoneLarge Matrix

#upload iphone LargeMatrix

```
iphone_LargeMatrix <- read.csv("iphone_LargeMatrix.csv.csv")
```

Data at a glance

```
summary(iphone_LargeMatrix)
```

```
str(iphone_LargeMatrix)
```

```
head(iphone_LargeMatrix)
```

```
names(iphone_LargeMatrix)
```

```
[1] "id" "iphone" "samsunggalaxy" "sonyxpria"
[5] "nokialumina" "htcphone" "ios" "googleandroid"
[9] "iphonecampos" "samsungcampos" "sonycampos" "nokiacampos"
[13] "htccampos" "iphonecamneg" "samsungcamneg" "sonycamneg"
[17] "nokiacamneg" "htccamneg" "iphonecamunc" "samsungcamunc"
[21] "sonycamunc" "nokiacamunc" "htccamunc" "iphonedispos"
[25] "samsungdispos" "sonydispos" "nokiadispos" "htcdispos"
[29] "iphonedisneg" "samsungdisneg" "sonydisneg" "nokiadisneg"
[33] "htcdisneg" "iphonedisunc" "samsungdisunc" "sonydisunc"
[37] "nokiadisunc" "htcdisunc" "iphoneperpos" "samsungperpos"
[41] "sonyperpos" "nokiaperpos" "htcperpos" "iphoneperneg"
[45] "samsungperneg" "sonyperneg" "nokiaperneg" "htcperneg"
[49] "iphoneperunc" "samsungperunc" "sonyperunc" "nokiaperunc"
[53] "htcperunc" "iosperpos" "googleperpos" "iosperneg"
[57] "googleperneg" "iosperunc" "googleperunc"
```

From the list of attribute names we can see that the first name "id" is new and is not consistent with the small matrix names.

remove "id" attribute from data frame

```
iphone_LargeMatrix$id <- NULL
```

convert iphone sentiment to a factor

```
iphone_LargeMatrix$iphonesentiment <- as.factor(iphone_LargeMatrix$iphonesentiment)
```

```
str(iphone_LargeMatrix)
```

Make predictions about iphone sentiment in the iphone_LargeMatrix

```
iphone_LM_C5.0_predictions <- predict(iphone_RC_C5.0model, iphone_LargeMatrix)
```

```
summary(iphone_LM_C5.0_predictions)
```

```
  1    2    3    4
7754 800 1297 7600
```

Work with large data matrix to make predictions regarding galaxy sentiment

Galaxy_LargeMatrix

#upload galaxy_LargeMatrix

```
galaxy_LargeMatrix<- read.csv("galaxy_LargeMatrix.csv.csv")
```

#Data at a glance

```
summary(galaxy_LargeMatrix)
```

```
str(galaxy_LargeMatrix)
```

```
names(galaxy_LargeMatrix)
```

#remove the "id" feature from the data set

```
galaxy_LargeMatrix$id <- NULL
```

#convert galaxy sentiment to factor

```
galaxy_LargeMatrix$galaxysentiment <- as.factor(galaxy_LargeMatrix$galaxysentiment)
```


```
str(galaxy_LargeMatrix)
```


Make predictions regarding galaxy sentiment on the galaxy Large Matrix

```
galaxy_LM_C5.0_predictions <- predict(galaxyRC_C5.0model, galaxy_LargeMatrix)
```

```
summary(galaxy_LM_C5.0_predictions)
```

```
1 2 3 4  
7814 792 1192 7653
```





The charts show a summary of the predicted classification of the iPhone and Galaxy smart phones listed in the large data matrix collected via Amazons EMR. Predictions were made based on the best fit machine learning algorithm, feature selection, and feature engineering approach. The large data matrix consists of over 17400 observations and 59 features. A. Show the actual count of each category. B. shows the percent of total count sentiment for both iphone and galaxy handsets with negative = 44% and 45 %, Neutral (somewhat neg & somewhat pos combined) = 12% and 11% and positive = 44% and 44% respectively.

In conclusion: it is clear that the sentiment for both iPhone and galaxy handsets are extremely similar. In making a choice for the medical app. One should consider the cost of the unit along with the availability of the unit in the countries of intended use.